

1. DCT Image Compression

```

Q1 > 1.py > ...
1 import numpy as np
2 import matplotlib
3 import matplotlib.pyplot as plt
4 import cv2
5 import math
6 import scipy
7
8 def imshow(image):
9     if len(image.shape) == 3:
10         # Height, width, channels
11         pass
12     else:
13         # Height, width - must be grayscale
14         # convert to RGB, since matplotlib will plot in a weird colormap (instead of black = 0, white = 1)
15         image = cv2.cvtColor(image, cv2.COLOR_GRAY2RGB)
16     # Draw the image
17     plt.imshow(image)
18     # Disable drawing the axes and tick marks in the plot, since it's actually an image
19     plt.axis('off')
20     plt.show()
21
22 n = int(input('Enter n:'))
23 m = int(input('Enter m:'))
24 select_img = int(input('Enter 1 for Barbara, 2 for Cat:'))
25 if select_img == 1:
26     input_img = cv2.imread('./Barbara.jpg')
27     img_name = 'bar'
28 elif select_img == 2:
29     input_img = cv2.imread('./cat.jpg')
30     img_name = 'cat'
31 input_img_rgb = cv2.cvtColor(input_img, cv2.COLOR_BGR2RGB)
32 # imshow(input_img_rgb)

```

- For this problem, first I import needed modules such as numpy, matplotlib, cv2, math, and scipy. I also initialize an imshow() function to show an image. Then I ask input for value of “n”, “m”, and to select the picture to be compressed. Then I read the selected image with imread() function, and convert it from BGR to RGB with cvtColor() function.

```

34 def keep_n_per_block(img, n):
35     x_lim = img.shape[0]
36     y_lim = img.shape[1]
37     result_img = np.zeros((x_lim, y_lim))
38
39     for j in range(0, y_lim, 8):
40         for i in range(0, x_lim, 8):
41             result_img[i:(i+n), j:(j+n)] = img[i:(i+n), j:(j+n)]
42     return result_img

```

- I also initialize a keep_n_per_block() function to keep only the lower-frequency (i.e. upper-left-n-by-n) coefficients in the 2D DCT domain 8x8 block, and set the rest of the coefficients in the block to 0.

(a) Implement the simplified DCT compression process for n = 2, 4 and m = 4, 8.

```

47 #COMPRESS
48 rgb_quantization_table = np.array([[ 8, 6, 6, 7, 6, 5, 8, 7],
49                                     [ 7, 7, 9, 9, 8,10,12,20],
50                                     [13,12,11,11,12,25,18,19],
51                                     [15,20,29,26,31,30,29,26],
52                                     [28,28,32,36,46,39,32,34],
53                                     [44,35,28,28,40,55,41,44],
54                                     [48,49,52,52,52,31,39,57],
55                                     [61,56,50,60,46,51,52,50]])
56
57
58 input_img_rgb_float = input_img_rgb.astype(np.float32)
59 x_lim = input_img_rgb.shape[0]
60 y_lim = input_img_rgb.shape[1]
61
62 #divide the img to 3 channels
63 input_img_r_float = np.zeros((x_lim, y_lim))
64 input_img_g_float = np.zeros((x_lim, y_lim))
65 input_img_b_float = np.zeros((x_lim, y_lim))
66 for j in range(0, y_lim):
67     for i in range(0, x_lim):
68         input_img_r_float[i, j] = input_img_rgb_float[i, j][0]
69         input_img_g_float[i, j] = input_img_rgb_float[i, j][1]
70         input_img_b_float[i, j] = input_img_rgb_float[i, j][2]
71

```

- (b) For part (a), first I create the rgb quantization table as 8x8 numpy array. Then I change the datatype of the input image to float, get the size of the image, and divide the image into 3 numpy arrays which are arrays for r, g, and b channels. I store the r channel to “input_img_r_rgb”, g channel to “input_img_g_rgb”, and b channel to “input_img_b_rgb”.

```

94 dct_img_r = np.zeros((x_lim, y_lim))
95 dct_img_g = np.zeros((x_lim, y_lim))
96 dct_img_b = np.zeros((x_lim, y_lim))
97
98 for j in range(0, y_lim, 8):
99     for i in range(0, x_lim, 8):
100         dct_img_r[i:(i+8), j:(j+8)] = dct2d(input_img_r_float[i:(i+8), j:(j+8)])
101         dct_img_g[i:(i+8), j:(j+8)] = dct2d(input_img_g_float[i:(i+8), j:(j+8)])
102         dct_img_b[i:(i+8), j:(j+8)] = dct2d(input_img_b_float[i:(i+8), j:(j+8)])
103
104 dct_img_r = keep_n_per_block(dct_img_r, n)
105 dct_img_g = keep_n_per_block(dct_img_g, n)
106 dct_img_b = keep_n_per_block(dct_img_b, n)
107

```

- I create a function to apply DCT and IDCT on 2D image called `dct2d()` and `idct2d()` function based on the formula on PPT. Then for each channel's image, I divide the image into blocks of 8x8 pixels and apply 2D DCT for each block using `dct2d()` function I define earlier, and store each channel's DCT coefficient to an numpy array.
- After it I keep only upper left nxn block coefficients in the 2D DCT domain for each 8x8 block by invoking `keep_n_per_block()` function, passing the result of DCT and n as arguments.

```

87 #Quantize using quantization table
88 for j in range(0, y_lim, 8):
89     for i in range(0, x_lim, 8):
90         for y in range(0, 8):
91             for x in range(0, 8):
92                 dct_img_r[(i+x), (j+y)] = round(dct_img_r[(i+x), (j+y)]/rgb_quantization_table[x, y])
93                 dct_img_g[(i+x), (j+y)] = round(dct_img_g[(i+x), (j+y)]/rgb_quantization_table[x, y])
94                 dct_img_b[(i+x), (j+y)] = round(dct_img_b[(i+x), (j+y)]/rgb_quantization_table[x, y])
95

```

- Then I quantize using rgb quantization table which I initialized earlier. I implement it using nested for loops, and divide each 8x8 block of the DCT coefficients with the 8x8 rgb quantization table.

```

103 #uniform quantization in m bits
104 max_value = max(np.max(dct_img_r), np.max(dct_img_g), np.max(dct_img_b))
105 min_value = min(np.min(dct_img_r), np.min(dct_img_g), np.min(dct_img_b))
106 step_size = (max_value - min_value) / 2**m
107 offset = round(min_value / step_size)
108 # print(max_value, min_value, step_size)
109
110 for j in range(0, y_lim):
111     for i in range(0, x_lim):
112         dct_img_r[i, j] = round(dct_img_r[i, j]/ step_size) - offset
113         dct_img_g[i, j] = round(dct_img_g[i, j]/ step_size) - offset
114         dct_img_b[i, j] = round(dct_img_b[i, j]/ step_size) - offset
115

```

- After quantization with quantization table, I do another quantization which is uniform quantization in m bits. First, I calculate the step size which is max value – min value of coefficients in each channel, divided by 2^m . I also calculate the offset value which are going to be added to the final value of the quantized coefficients to make all of the coefficients to be non-negative values and easier to save as an image. Note that this offset will not increase the number of bits used, it will just shift the quantized coefficients to the positive side. I implement this uniform quantization by nested for loops, and I divide each DCT coefficients by the step size, and subtract it by the offset value.

```

119 #combine 3 channels to 1 img
120 dct_img_rgb = np.zeros((x_lim, y_lim, 3))
121 # dct_img_rgb = input_img_rgb.copy() #just to get array with same size
122 for j in range(0, y_lim):
123     for i in range(0, x_lim):
124         dct_img_rgb[i, j][0] = dct_img_r[i, j]
125         dct_img_rgb[i, j][1] = dct_img_g[i, j]
126         dct_img_rgb[i, j][2] = dct_img_b[i, j]
127

```

- Finally, I combine all the 3 channel's DCT coefficients to an image by using nested for loops. First I initialize an numpy array with the same size as input image, and for each pixel, I assign each channel's DCT coefficients to be the value of this image. Then we are done for the DCT compression.

```

129 #DECOMPRESS
130
131 #divide the img to 3 channels
132 idct_img_r = np.zeros((x_lim, y_lim))
133 idct_img_g = np.zeros((x_lim, y_lim))
134 idct_img_b = np.zeros((x_lim, y_lim))
135 for j in range(0, y_lim):
136     for i in range(0, x_lim):
137         idct_img_r[i, j] = dct_img_rgb[i, j][0]
138         idct_img_g[i, j] = dct_img_rgb[i, j][1]
139         idct_img_b[i, j] = dct_img_rgb[i, j][2]
140

```

- Before I start to decompress the compressed image, I first divide the compressed image to 3 channels using nested for loops, the separated results are "idct_img_r" for red channel, "idct_img_g" for green channel, and "idct_img_b" for blue channel. These values are still the quantized dct coefficients, I call this idct only to differentiate the compression and decompression process, and because these values are going to be applied idct later.

```

141 #Uniform Unquantization
142 for j in range(0, y_lim):
143     for i in range(0, x_lim):
144         idct_img_r[i, j] = (idct_img_r[i, j] + offset)*step_size
145         idct_img_g[i, j] = (idct_img_g[i, j] + offset)*step_size
146         idct_img_b[i, j] = (idct_img_b[i, j] + offset)*step_size
147

```

- Then I unquantize the dct coefficients with uniform unquantization. I implement this by using nested for loops, and for each pixel, I add the values by offset values, and multiply it by step size.

```

148 #Unquantize using quantization table
149 for j in range(0, y_lim, 8):
150     for i in range(0, x_lim, 8):
151         for y in range(0, 8):
152             for x in range(0, 8):
153                 idct_img_r[(i+x), (j+y)] = idct_img_r[(i+x), (j+y)]*rgb_quantization_table[x, y]
154                 idct_img_g[(i+x), (j+y)] = idct_img_g[(i+x), (j+y)]*rgb_quantization_table[x, y]
155                 idct_img_b[(i+x), (j+y)] = idct_img_b[(i+x), (j+y)]*rgb_quantization_table[x, y]
156

```

- Then I unquantize the coefficients again with quantization table. I implement it using nested for loops, and multiply each 8x8 block of these coefficients with the 8x8 rgb quantization table.

```

195 #IDCT for each channel
196 result_img_r = np.zeros((x_lim, y_lim))
197 result_img_g = np.zeros((x_lim, y_lim))
198 result_img_b = np.zeros((x_lim, y_lim))
199
200 for j in range(0, y_lim, 8):
201     for i in range(0, x_lim, 8):
202         result_img_r[i:(i+8), j:(j+8)] = idct2d(idct_img_r[i:(i+8), j:(j+8)])
203         result_img_g[i:(i+8), j:(j+8)] = idct2d(idct_img_g[i:(i+8), j:(j+8)])
204         result_img_b[i:(i+8), j:(j+8)] = idct2d(idct_img_b[i:(i+8), j:(j+8)])
205

```

- Then I apply IDCT to these DCT coefficients to get the decompressed value of the image. For each channel's DCT coefficients, I divide the coefficients into blocks of 8x8 pixels and apply 2D IDCT for each block using idct2d() function, and store each channel's resulting value to a numpy array.

```

168 #combine 3 channels to 1 img
169 result_img_rgb = input_img_rgb.copy() #just to make array with same size
170 for j in range(0, y_lim):
171     for i in range(0, x_lim):
172         R = result_img_r[i, j]
173         G = result_img_g[i, j]
174         B = result_img_b[i, j]
175         #prevent overflow or underflow
176         result_img_rgb[i, j][0] = R if R <= 255 and R >= 0 else 255 if R > 255 else 0 #R part
177         result_img_rgb[i, j][1] = G if G <= 255 and G >= 0 else 255 if G > 255 else 0 #G part
178         result_img_rgb[i, j][2] = B if B <= 255 and B >= 0 else 255 if B > 255 else 0 #B part
179

```


- Then I combine the 3 arrays that store the resulting value of idct to 1 image. I implement this by using nested for loops, and I set 255 as the upper limit and 0 as the under limit of the values to prevent overflow and underflow.

```


180 imshow(result_img_rgb)
181 cv2.imwrite('./output/'+img_name+'_'+n+'m'+str(m)+'_a.jpg', cv2.cvtColor(result_img_rgb, cv2.COLOR_RGB2BGR))
182
183 #compute compression ratios
184 compress_ratio = (x_lim*y_lim*8*3)/(x_lim*y_lim*m)
185 print('Compression ratio:', compress_ratio)
186
187 #compute PSNR values
188 mse = np.mean((input_img_rgb.astype(np.float32)-result_img_rgb.astype(np.float32))**2)
189 max_pixel = 255.0
190 psnr = 20*math.log10(max_pixel / math.sqrt(mse))
191 print('PSNR value:', psnr)

```

- Then I show the decompressed image, and convert it from RGB to BGR color channel before writing the image to a jpg file.
- And finally, I compute the compression ratios and PSNR values of the compressed image. Since in the original image each pixel uses 24 bits (8 bits x 3 channels), and in the compressed image each pixel only used m bits, so the compression ratio is 24/m. For PSNR values, I calculate it by first calculate the mean squared error value, and since the peak value of a pixel is 255, the PSNR is 20 multiplied by log 10 of (255/mse).



Original




n=2, m=4



n=2, m=8

Compression ratio: 6.0
PSNR value: 24.183193610466148

Compression ratio: 3.0
PSNR value: 25.612204866470023




n=4, m=4



n=4, m=8

Compression ratio: 6.0
PSNR value: 24.802319105392108

Compression ratio: 3.0
PSNR value: 28.996261616913316




Original




$n=2, m=4$

Compression ratio: 6.0
PSNR value: 22.869370524163557




$n=2, m=8$

Compression ratio: 3.0
PSNR value: 24.110154650606578



$n=4, m=4$

Compression ratio: 6.0
PSNR value: 25.60261390794929



$n=4, m=8$

Compression ratio: 3.0
PSNR value: 33.16800271987338

Discussion for (a)

- When $m=4$, the compression ratio is 6.0, which causes the decompressed image to look really grainy. But for $m=8$, although the compression ratio is less with 3.0, but the decompressed image look more clearer and the difference with the original image is not that obvious.
- For $m=4$, the difference from changing $n=2$ to $n=4$ is not much for Barbara image, but a lot for cat image. For $m=8$, the difference from changing $n=2$ to $n=4$ is a lot for both Barbara image and cat image.
- For $n=2$, PSNR value of Barbara image is larger than PSNR value of cat image, meaning the distortion rate of the compressed Barbara image is less than the distortion rate of the compressed cat image. But for $n=4$, PSNR value of Barbara image is smaller than the cat image. Meaning the compression with $n=2$ works better for Barbara image and compression with $n=4$ works better for cat image.

(c) Use the same process in (a) with image transformed to YCbCr color space with 4:2:0 chrominance subsampling.

```

196 #(b)
197
198 input_img_ycc = input_img_rgb.copy()
199 x_lim = input_img_ycc.shape[0]
200 y_lim = input_img_ycc.shape[1]
201
202 #convert input img from RGB color space to YCbCr color space
203 for j in range(0, y_lim):
204     for i in range(0, x_lim):
205         input_img_ycc[i, j][0] = 16+ 0.257*input_img_rgb[i, j][0] + 0.564*input_img_rgb[i, j][1] + 0.098*input_img_rgb[i, j][2] #Y part
206         if j%2 == 0 and i%2 == 0: #sample chrominance
207             Cb = 128 - 0.148*input_img_rgb[i, j][0] - 0.291*input_img_rgb[i, j][1] + 0.439*input_img_rgb[i, j][2]
208             Cr = 128 + 0.439*input_img_rgb[i, j][0] - 0.368*input_img_rgb[i, j][1] - 0.071*input_img_rgb[i, j][2]
209             #4:2:0 chrominance subsampling
210             for y in range(j, j+2):
211                 for x in range(i, i+2):
212                     if y < y_lim and x < x_lim:
213                         input_img_ycc[x, y][1] = Cb
214                         input_img_ycc[x, y][2] = Cr
215

```

- For part (b), I first convert the image from RGB color space to YCbCr color space with 4:2:0 chrominance subsampling. I implement this by using nested for loops, and for every pixel I calculate the Y value of the pixel with the formula provided by the TA in discussion forum. And if the indexes of the pixel in x-axis and y-axis are both multiply of 2, I calculate the chrominance value by the formula, and use this chrominance value for the pixel, the pixel's right pixel, the pixel under, and the pixel on right under side of this pixel.

```

217 #COMPRESS
218 lumin_quantization_table = np.array([[16,11,10,16,24,40,51,61],
219 [12,12,14,19,26,58,60,55],
220 [14,13,16,24,40,57,69,56],
221 [14,17,22,29,51,87,80,62],
222 [18,22,37,56,68,109,103,77],
223 [24,36,55,64,81,104,113,92],
224 [49,64,78,87,103,121,120,101],
225 [72,92,95,98,112,100,103,99]])
226
227 chrom_quantization_table = np.array([[17,18,24,47,99,99,99,99],
228 [18,21,26,66,99,99,99,99],
229 [24,26,56,99,99,99,99,99],
230 [47,66,99,99,99,99,99,99],
231 [99,99,99,99,99,99,99,99],
232 [99,99,99,99,99,99,99,99],
233 [99,99,99,99,99,99,99,99],
234 [99,99,99,99,99,99,99,99]])
235

```

- Then I initialize the luminance quantization table and chrominance quantization table.
- After it I compress the image in YCbCr color space the similarly as in (a), the only differences are when I quantize a value using quantization table, instead of using rgb quantization table, I use these luminance and chrominance quantization table. And another difference is for the uniform quantization, because the value of the luminance part and chrominance part differs quite a lot, I separate the uniform quantization to 2 part, one for the luminance part and one for the chrominance part. Note that separating the uniform quantization won't increase the number of bits required for each pixel, as both part are using the same range of numbers.

```

340 #IDCT for each channel
341 result_img_y = np.zeros((x_lim, y_lim))
342 result_img_cb = np.zeros((x_lim, y_lim))
343 result_img_cr = np.zeros((x_lim, y_lim))
344 for j in range(0, y_lim, 8):
345     for i in range(0, x_lim, 8):
346         result_img_y[i:(i+8), j:(j+8)] = cv2.idct(idct_img_y[i:(i+8), j:(j+8)])
347         result_img_cb[i:(i+8), j:(j+8)] = cv2.idct(idct_img_cb[i:(i+8), j:(j+8)])
348         result_img_cr[i:(i+8), j:(j+8)] = cv2.idct(idct_img_cr[i:(i+8), j:(j+8)])
349
350
351
352 #convert img from YCbCr color space to RGB color space
353 converted_img = input_img_rgb.copy() #just to make array with same shape
354
355 for j in range(0, y_lim):
356     for i in range(0, x_lim):
357         R = 1.164*(result_img_y[i, j]-16) + 0*(result_img_cb[i, j]-128) + 1.596*(result_img_cr[i, j]-128)
358         G = 1.164*(result_img_y[i, j]-16) - 0.382*(result_img_cb[i, j]-128) - 0.813*(result_img_cr[i, j]-128)
359         B = 1.164*(result_img_y[i, j]-16) + 2.017*(result_img_cb[i, j]-128) + 0*(result_img_cr[i, j]-128)
360         #prevent overflow or underflow
361         converted_img[i, j][0] = R if R <= 255 and R >= 0 else 255 if R > 255 else 0 #R part
362         converted_img[i, j][1] = G if G <= 255 and G >= 0 else 255 if G > 255 else 0 #G part
363         converted_img[i, j][2] = B if B <= 255 and B >= 0 else 255 if B > 255 else 0 #B part

```


- After applying IDCT to each channel, I convert the image from YCbCr color space back to RGB color space. I implement this using nested for loops. For each pixel I calculate the R, G, and B values from Y, Cb, Cr value using the formula, and with similar ways as in part (a) I prevent overflow and underflow when assigning the values of R, G, and B to the converted image.

```

365 imshow(converted_img)
366 cv2.imwrite('./output/'+img_name+'_n'+str(n)+'m'+str(m)+'_b.jpg', cv2.cvtColor(converted_img, cv2.COLOR_RGB2BGR))
367
368 #compute compression ratios
369 compress_ratio = (x_lim*y_lim*8*3)/(x_lim*y_lim*m*2)
370 print('Compression ratio:', compress_ratio)
371
372 #compute PSNR values
373 mse = np.mean((input_img_rgb.astype("float")-converted_img.astype("float"))**2)
374 max_pixel = 255.0
375 psnr = 20*math.log10(max_pixel / math.sqrt(mse))
376 print('PSNR value:', psnr)

```


- Finally, I show the image and convert it to BGR color space before writing it as jpg file, and I calculate the compression ratio and PSNR values of the compressed image.



Original



n=2, m=4




n=2, m=8

<p>Compression ratio: 6.0</p> <p>PSNR value: 22.867156239363812</p>	<p>Compression ratio: 3.0</p> <p>PSNR value: 24.59709866216589</p>
---	--




n=4, m=4

Compression ratio: 6.0
PSNR value: 23.468393955946013



n=4, m=8

Compression ratio: 3.0
PSNR value: 26.906162649719263




Original



n=2, m=4

Compression ratio: 6.0
PSNR value: 21.980021548395428




n=2, m=8

Compression ratio: 3.0
PSNR value: 23.133265577842977



n=4, m=4

Compression ratio: 6.0
PSNR value: 24.29434131681066



n=4, m=8

Compression ratio: 3.0
PSNR value: 28.205142267360635

Discussion for (b)

- When $m=4$, the compression ratio is 6.0, which causes the decompressed image to look really grainy. But for $m=8$, although the compression ratio is less with 3, but the decompressed image looks more clearer and the difference with the original image is not that obvious.
- For $m=4$, the difference from changing $n=2$ to $n=4$ is not much for Barbara image, but a lot for cat image. For $m=8$, the difference from changing $n=2$ to $n=4$ is a lot for both Barbara image and cat image.
- For $n=2$, PSNR value of Barbara image is larger than PSNR value of cat image, meaning the distortion rate of the compressed Barbara image is less than the distortion rate of the compressed cat image. But for $n=4$, PSNR value of Barbara image is smaller than the cat image. Meaning the compression with $n=2$ works better for Barbara image and compression with $n=4$ works better for cat image.

Discussion

- The decompressed image with compression on YCbCr color spaces are more blocky when $m=4$, and a little bit more yellowish than the decompressed image with compression on RGB color spaces.
- I think the compression quality with RGB color space is better than compression on YCbCr color space as the PSNR values of images compressed in RGB color space are larger than PSNR values of images compressed in YCbCr color space.

2. Create your own FIR filters to filter audio signal

```
Q2 > 2.py > ...
1  import numpy as np
2  import matplotlib
3  import matplotlib.pyplot as plt
4  import cv2
5  import math
6  import scipy.io.wavfile
7
8  audio = scipy.io.wavfile.read("./HW2_Mix_2.wav")
9  sampling_rate = audio[0]
10 print('Audio sampling rate:', sampling_rate)
11 audio = np.array(audio[1], dtype=float)
12
13 #TIME DOMAIN SIGNAL
14 plt.plot(audio[:])
15 plt.ylabel("Amplitude")
16 plt.xlabel("Time")
17 plt.show()
```

- Before starting, first I import needed modules such as numpy, matplotlib, math, and scipy.io.wavfile. Then I read the audio to be processed with scipy.io.wavfile.read() function, the return value of this function is an array with 2 elements, the first element is the sampling rate and the second element is the signal data of the audio.


- Then I using matplotlib module, I plot the signal data of the audio and show it.

```

19 #FREQUENCY DOMAIN SIGNAL
20 # Calculate n/2 to normalize the FFT output
21 n = audio.size
22 normalize = n/2
23 fft_signal = np.fft.fft(audio)
24 fft_freq = np.fft.fftfreq(n, d=1.0/sampling_rate)
25
26 plt.plot(fft_freq, np.abs(fft_signal/normalize))
27 plt.ylabel("Amplitude")
28 plt.xlabel("Frequency")
29 plt.xlim(-2000, 2000)
30 plt.savefig('./output/input.png')
31 plt.show()

```

- I transform the time domain signal to Frequency domain signal using Fast Fourier Transform. I implement fast fourier transform by using `np.fft.fft()` and `np.fft.fftfreq()` functions to get the values on x-axis and y-axis, and I then plot them and save the plot as "input.png". Here before saving I set the x-axis min limit to -2000 and max limit to 2000 to make the signal larger and clearer to see.



- Based on the frequency domain graph, I notice that the frequencies is divided into 3 big parts, around 0Hz-380Hz, 380Hz-780Hz, and 780Hz-1550Hz. So, based on this observation I decided to use lowpass filter as the first filter to filter frequencies ranging in 0Hz-380Hz, bandpass filter as the second filter for frequencies ranging around 380Hz-780Hz, and highpass filter as the third filter for frequencies at and larger than 780Hz.

```

38 # FILTER 1 (LOWPASS FILTER)
39 cutoff_freq = 380
40 cutoff_freq = cutoff_freq/sampling_rate
41 w_c = 2*math.pi*cutoff_freq
42 filter_order = 32000
43 middle = filter_order//2
44 low_filter = np.zeros((filter_order), dtype=float)
45
46 for i in range(-middle, middle):
47     if i == 0:
48         low_filter[middle] = 1
49     else:
50         low_filter[i+middle] = math.sin(2*math.pi*cutoff_freq*i)/(math.pi*i)
51 low_filter[middle] = 2*cutoff_freq
52
53 #multiply the elements of filter by a windowing function
54 #blackmann window function
55 for i in range(0, filter_order):
56     low_filter[i] = low_filter[i] * (0.42 + 0.5*math.cos((2*math.pi*i)/(n-1)) + 0.08*math.cos((4*math.pi*i)/(n-1)))

```

- For the first filter (lowpass filter), I implement it by first designing the filter. I set the cutoff frequency and the filter order of the filter, and initialize a numpy array to store the coefficients of the filter. Then I use for loops to set the filter coefficient values using the formula $\sin(2\pi \cdot \text{cutoff_freq} \cdot i) / (\pi \cdot i)$.
- After the for loops is done, I again use another for loop to multiply the element of filter by a blackmann windowing function. For each element in the filter, I multiply it with the blackmann windowing formula.

```

61 #1D CONVOLUTION
62 filtered_signal_1 = audio.copy() #just to get array with same size
63 # Pad the input signal with zeros to compute the convolution in the "same" mode
64 pad_length = len(low_filter) // 2
65 audio_padded = np.pad(audio, (pad_length, pad_length), mode='constant')
66 # Compute the convolution
67 for i in tqdm.trange(len(audio)):
68     filtered_signal_1[i] = np.dot(audio_padded[i:i+len(low_filter)], low_filter)
69

```

- After the filter is completely initialized, I start the 1D convolution. The idea to implement it is to first pad the input signal with zeros so that the output signal has the same length as the input signal, then I perform the convolution. In the code, I first compute the padding length as half the length of the filter, then I pad the input signal with zeros using the np.pad(), and I compute the convolution by iterating over the indices of the output signal and using the np.dot() function to compute the dot product of the corresponding slice of the padded input signal and the filter.

```

70 #output audio
71 scipy.io.wavfile.write('./output/Filter1Lowpass_380.wav', sampling_rate, np.int16(filtered_signal_1))
72
73 #output spectrums
74 # Calculate n/2 to normalize the FFT output
75 n = filtered_signal_1.size
76 normalize = n/2
77 fft_signal_1 = np.fft.fft(filtered_signal_1)
78 fft_freq_1 = np.fft.fftfreq(n, d=1.0/sampling_rate) #(?)
79
80 plt.plot(fft_freq_1, np.abs(fft_signal_1/normalize))
81 plt.ylabel("Amplitude")
82 plt.xlabel("Frequency")
83 plt.title('Output spectrum by filter 1 (Lowpass)')
84 plt.xlim(-2000, 2000)
85 plt.savefig('./output/output_by_Filter1Lowpass.png')
86 plt.show()
87

```

- After having the output signal saved in “filtered_signal_1” array, I write the output signal using scipy.io.wavfile.write() function.

- I also transform the output signal to its frequency domain signal with similar ways as when I transform the input signals, and I save the figure before showing it.

```

88     #filter shape
89     plt.plot(low_filter)
90     plt.title('Filter 1 (Lowpass) shape')
91     plt.xlim(11000, 21000)
92     plt.savefig('./output/Filter1Lowpass_shape.png')
93     plt.show()
94
95     #filter spectrums
96     n = filter_order
97     normalize = n/2
98     fft_signal_1 = np.fft.fft(low_filter)
99     fft_freq_1 = np.fft.fftfreq(n, d=1.0/sampling_rate) #(?)
100    plt.plot(fft_freq_1, np.abs(fft_signal_1/normalize))
101    plt.title('Filter 1 (Lowpass) filter spectrum')
102    plt.ylabel("Amplitude")
103    plt.xlabel("Frequency")
104    plt.xlim(-2000, 2000)
105    plt.savefig('./output/Filter1Lowpass_spectrum.png')
106    plt.show()

```

- I also plot the shape of the filter in time domain, and set the x-axis limit to 11000 until 21000 to make it clearer, and save the figure before showing it.
- Then I transform the filter to its frequency domain with fast fourier transform, set the x-axis limit to -2000 until 2000, and save the figure before showing it.

```

109    #DOWN SAMPLE 2KHZ
110    original_sampling_rate = sampling_rate
111    new_sampling_rate = 2000
112    downsample_factor = original_sampling_rate / new_sampling_rate
113    num_samples_downsampled = int(len(filtered_signal_1)/downsample_factor)
114    downsampled_signal = np.zeros((num_samples_downsampled), dtype=float)
115    #downsample the signal
116    for i in range(num_samples_downsampled):
117        downsampled_signal[i] = filtered_signal_1[int(i * downsample_factor)]
118    #output audio
119    scipy.io.wavfile.write('./output/Filter1Lowpass_380_2khz.wav', new_sampling_rate, np.int16(downsampled_signal))

```

- To down sample the filtered signal to 2000Hz, I first set the audio sampling rate as the original sampling rate and 2000 as the new sampling rate. Then I calculate the down sample factor by dividing original sampling rate by new sampling rate.
- After it I calculate the number of samples there will be after the down sample, I calculate this by dividing the length of the filtered audio by the down sample factor. With this number I create a new numpy array to later store the down sampled signal.
- Then using for loop, I set the i-th index value of down sample signal to be the value of the (i*downsample factor)-th index value of the filtered signal.
- Then I write the down sampled signal as a wav file.

```

#ONE-FOLD ECHO
one_echo_filter = np.zeros((3201), dtype=float)
one_echo_filter[0] = 1
one_echo_filter[3200] = 0.8

one_echo_filtered_signal = filtered_signal_1.copy() #just to get array with same size
# Pad the input signal with zeros
pad_length = len(one_echo_filter) // 2
audio_padded = np.pad(filtered_signal_1, (pad_length, pad_length), mode='constant')
# Compute the convolution
for i in tqdm.trange(len(audio)):
    one_echo_filtered_signal[i] = np.dot(audio_padded[i:i+len(one_echo_filter)], one_echo_filter)
#output audio
scipy.io.wavfile.write('./output/Echo_one.wav', sampling_rate, np.int16(one_echo_filtered_signal))

```


- I also apply one-fold echo and multiple-fold echo to the filtered signal before downsampling. I implement one-fold echo by first initializing the one-echo filter, I do this by creating a numpy array of zeros with length 3201, and float as the data type. Then I set the first element of this filter to be 1 and last element of this filter to be 0.8.
- Then using the same way as the lowpass filter convolution earlier, I compute the one-echo filtered signal. The only difference is the filter I use for the convolution is one-echo filter instead of lowpass filter, and the filtered signal is the low pass filtered signal instead of the original audio.
- Then I write the one-echo filtered signal as wavfile.

```

137 #MULTIPLE-FOLD ECHO
138 multi_echo_filter = np.zeros((3201), dtype=float)
139 multi_echo_filter[0] = 1
140 multi_echo_filter[3200] = -0.8
141
142 multi_echo_filtered_signal = filtered_signal_1.copy() #just to get array with same size
143 # Pad the input signal with zeros
144 pad_length = len(multi_echo_filter) // 2
145 audio_padded = np.pad(filtered_signal_1, (pad_length, pad_length), mode='constant')
146 # Compute the convolution
147 for i in tqdm.trange(len(audio)):
148     multi_echo_filtered_signal[i] = np.dot(audio_padded[i:i+len(multi_echo_filter)], multi_echo_filter)
149 #output audio
150 scipy.io.wavfile.write('./output/Echo_multiple.wav', sampling_rate, np.int16(multi_echo_filtered_signal))

```

- To apply the multiple-fold echo to the filtered signal before down sampling, I implement it the as how I implement the one-fold echo, the only difference is the filter. For the multi-echo filter, the last element value is - 0.8 instead of 0.8 in one-fold echo.

```

154 # FILTER 2 (BANDPASS FILTER)
155 freq_1 = 380
156 freq_2 = 780
157 freq_1 = freq_1/sampling_rate
158 freq_2 = freq_2/sampling_rate
159 filter_order = 32000
160 middle = filter_order//2
161 bandpass_filter_2 = np.zeros((filter_order), dtype=float)
162
163 for i in range(-middle, middle):
164     if i == 0:
165         bandpass_filter_2[middle] = 1
166     else:
167         bandpass_filter_2[i+middle] = math.sin(2*math.pi*freq_2*i)/(math.pi*i) - math.sin(2*math.pi*freq_1*i)/(math.pi*i)
168 bandpass_filter_2[middle] = 2*(freq_2-freq_1)
169
170 #blackmann window function
171 for i in range(0, filter_order):
172     bandpass_filter_2[i] = bandpass_filter_2[i] * (0.42 + 0.5*math.cos((2*math.pi*i)/(n-1)) + 0.08*math.cos((4*math.pi*i)/(n-1)))
173

```

- For the second filter, which is bandpass filter for frequencies ranging in 380Hz-780Hz, I implement it by first initializing a numpy array with filter order as its length to store the coefficients of this bandpass filter. I also set the 380 as the frequency 1 and 780 as frequency 2.
- Then I use nested for loops to set the filter coefficient values using the formula $\sin(2\pi \cdot \text{frequency } 1 \cdot i) / (\pi \cdot i)$ minus $\sin(2\pi \cdot \text{frequency } 2 \cdot i) / (\pi \cdot i)$.
- After the for loops is done, I again use another for loop to multiply the element of filter by a blackmann windowing function. For each element in the filter, I multiply it with the blackmann windowing formula.

```

174 filtered_signal_2 = audio.copy() #just to get array with same size
175 # Pad the input signal with zeros
176 pad_length = len(bandpass_filter_2) // 2
177 audio_padded = np.pad(audio, (pad_length, pad_length), mode='constant')
178 # Compute the convolution
179 for i in tqdm.trange(len(audio)):
180     filtered_signal_2[i] = np.dot(audio_padded[i:i+len(bandpass_filter_2)], bandpass_filter_2)

```

- Then I compute the convolution of the original audio and the filter, save it to “filtered_signal_2” array, and write it as wav file.
- Then I show the output spectrums, filter shape, and filter spectrums and save it as png file with the same way I do it in the first filter.
- I also down sample this bandpass filtered signal to 2000Hz, and write it as wav file. I implement this with the same way when I down sample the lowpass filtered signal to 2000Hz.

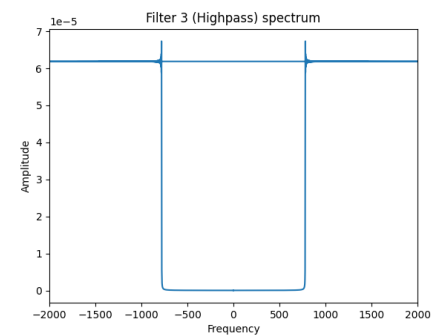
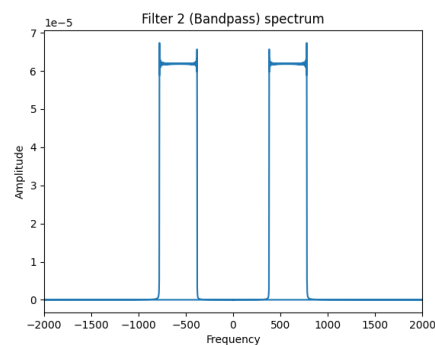
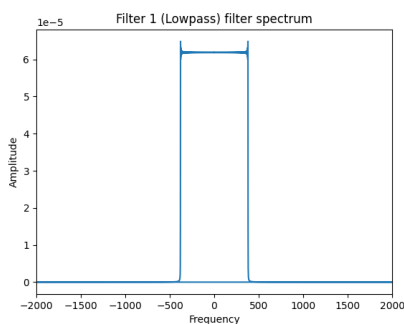
```

234 #FILTER 3 (HIGHPASS FILTER)
235 cutoff_freq = 780
236 cutoff_freq = cutoff_freq/sampling_rate
237 w_c = 2*math.pi*cutoff_freq
238 filter_order = 32000
239 middle = filter_order//2
240 highpass_filter = np.zeros((filter_order), dtype=float)
241
242 for i in range(-middle, middle):
243     if i == 0:
244         highpass_filter[middle] = 1
245     else:
246         highpass_filter[i+middle] = -math.sin(2*math.pi*cutoff_freq*i)/(math.pi*i)
247     highpass_filter[middle] = 1 - 2*cutoff_freq

```


- As for the 3rd filter which is a highpass filter, I implement it similarly as the way I implement the first filter. The only difference is the value of the cutoff frequency which is 780 here instead of 380, and the formula to calculate the filter coefficient is changed to $-\sin(2\pi \cdot \text{cutoff_freq} \cdot i) / (\pi \cdot i)$, and $1 - 2 \cdot \text{cutoff_freq}$ for when $i = 0$.
- For the other part it is all the same as filter 1.
- I also down sample this bandpass filtered signal the same way as I down sample filtered signal in other filters.

Discussion



- The spectrum of the second filter is separated to two blocks like the 2 other filters as the second filter filters frequencies ranging in the middle (380Hz-780Hz).

- The total width of a block in spectrum of the first and second filter is narrower than the total width of a block in spectrum of the third filter. This is because the range of frequencies that the third filter filters on (780Hz-22500Hz) is larger than the other 2 filters.



- The higher the frequencies a filter is filtering on, the tighter the filter shape. I think it is because higher frequencies will have more waves density, which causes the shape of the filter to be tighter.

Briefly compare the difference between signals before and after reducing the sampling rates.

- For the signals filtered by first filter and second filter which are signals with frequencies ranging around 0Hz-380Hz and 380Hz-780Hz respectively, down sampling to 2000Hz have no significant difference to our ears, as according to Nyquist's sampling theorem, 2000Hz is more than twice the highest frequency component in the signals filtered by first filter and second filter.
- But for signals filtered by third filter, which are signals with frequencies ranging around 780Hz-1550Hz, down sampling to 2000Hz will results in aliasing as 2000Hz is less than twice the highest frequency component in the signals filtered, and there aren't enough sample points from which to accurately interpolate the sinusoidal form of the filtered wave.